

ALMA MATER STUDIORUM · UNIVERSITÀ DI BOLOGNA

DIPARTIMENTO DI FISICA E ASTRONOMIA - DIFA

PhD Thesis seminary

---

# Implementation and optimization of algorithms in Biomedical Big Data Analytics

Nico Curti

**Supervisor:**

*Prof. Daniel Remondini*

**Co-Supervisor:**

*Prof. Gastone Castellani*

*Prof. Armando Bazzani*

January 27, 2020



# Purposes

## Biomedical Big Data Analytics

### Feature Selection DNetPRO

- ◊ **Topic:** Gene Expression Analysis;
- ◊ **Characteristics:** Novel algorithm & optimization on distributed computing;
- ◊ **Applications:**
  - ▶ Synthetic dataset;
  - ▶ Benchmark TCGA datasets;
  - ▶ Cytokinoma;
  - ▶ Bovine Signature;
  - ▶ Pedestrian mobility in Venice;
- ◊ **URL:** <https://github.com/Nico-Curti/DNetPRO>;

### Deep Learning Byron

- ◊ **Topic:** Neural Network models;
- ◊ **Characteristics:** Code Optimization and multi-threading parallelization;
- ◊ **Applications:**
  - ▶ Super Resolution;
  - ▶ Object Detection;
  - ▶ Image Segmentation;
- ◊ **URL:**
  - ▶ <https://github.com/Nico-Curti/NumPyNet>;
  - ▶ <https://github.com/Nico-Curti/Byron>;

### Big Data CHIMeRA

- ◊ **Topic:**
  - ▶ Biomedical Big Data;
  - ▶ Database service;
  - ▶ Network analysis;
  - ▶ Natural Language Processing;
  - ▶ Web scraping;
  - ▶ Data harmonization;
- ◊ **Characteristics:** Merging of public available knowledge into network structure;
- ◊ **Applications:** Leukemia;
- ◊ **URL:** Not available yet;



# Deep Learning

## Neural Network models

Deep Learning model is becoming equivalent to Neural Network architecture/model, i.e a more or less complex pipeline of functions which takes in input a sample and it applies a series of transformations to obtain the desired result.

### Objectives

- Optimization and extension of state-of-art Neural Network libraries;
- Development of two novel libraries, for education and analytic purposes, respectively;
- Improve parallelization strategies to reach the best computational performances in cluster machine without GPUs supports;

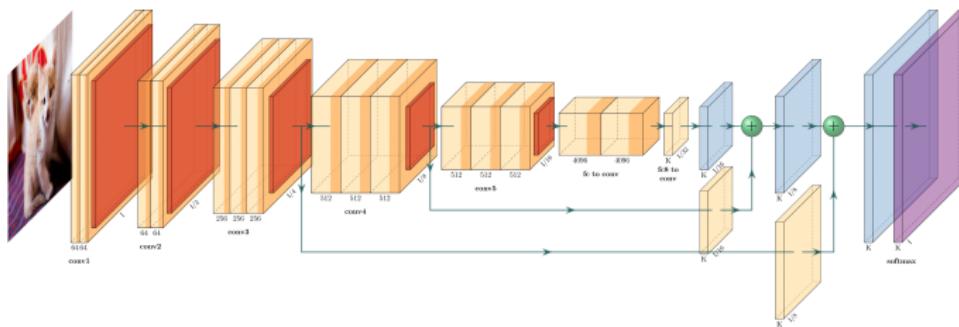
All the algorithm and codes developed are **tested, documented and public available!**

**Reference:** <https://nico-curti2.gitbook.io/phd-thesis> (GitBook documentation).

# What is a Deep Learning model?

## Neural Network models

- Neural Networks are mathematical models **commonly used** in data analysis.
- From a theoretical point-of-view we can define a Neural Network as **a series of non-linear multi-parametric functions**.
- Neural Networks are considered as **Universal Function Approximators**.



**Figure:** Neural Network model scheme. The model is composed by a series of layers made by interconnected computational units. The connections between layers can be serials or parallels and each layer models a pre-determined mathematical function



# State-of-art libraries

## Why other two libraries?

A wide range of documentations and implementations have been written on Deep Learning and it is more and more hard to move around the different sources. Leader on this topic are became the multiple open-source ‘Python’ libraries available on-line as [Tensorflow](#), [Pytorch](#) and [Caffe](#).

### Pros

- ◊ Simplicity in writing complex models in a minimum number of code lines;
- ◊ Simplicity of the Python language;
- ◊ Extremely efficient in GPU environments;

### Cons

- ◊ Steep learning curve for complex applications or algorithm modifications;
- ◊ Hard to manage as services;
- ◊ They are not designed for CPU performances;

Only a small part of the research community uses more deeper implementation in C++ or other low-level programming languages. About them it should be mentioned the [darknet](#) project of Redmon J. et al. which created a sort of standard in object detection applications using a pure Ansi-C library.



# NumPyNet & Byron

Why other two libraries?

## NumPyNet



- **LANGUAGE:** Python;
- **INSTALLATION:** Setup ([PyPI](#) upload wip);
- **OS:** Linux & MacOS & Windows;
- **DOCUMENTATION:** fully documented;
- **PURPOSE:** educational / study;
- **CI:** Travis & Appveyor;
- **USAGE:** model design / algorithm improvements;
- **MISCELLANEOUS:** all functionality are tested against **Keras** counterpart;
- **URL:** <https://github.com/Nico-Curti/NumPyNet>

## Byron



- **LANGUAGE:** C++ Standard 17;
- **INSTALLATION:** CMake & Make & Docker;
- **OS:** Linux & MacOS & Windows;
- **DOCUMENTATION:** wide set of usage examples;
- **PURPOSE:** multi-threading CPU optimization;
- **CI:** Travis & Appveyor & CircleCI;
- **USAGE:** Super Resolution & Object Detection & Image Segmentation;
- **MISCELLANEOUS:** extension and optimization of the **darknet** project;
- **URL:** <https://github.com/Nico-Curti/Byron>

From the union of this two project **Pyron** is born, i.e the completely wrap of the **Byron** library in **Python** using **Cython**.

\* Both libraries are already used in different works of thesis and research projects!



# NumPyNet

## Neural Networks in Pure NumPy

### NumPyNet



- ▶ The library is written in pure **Python** and the only **external library** used is **Numpy** (a based package for the scientific research).
- ▶ Despite all common libraries are correlated by a wide documentation is often difficult for novel users to move around the many hyper-links and papers cite in it. NumPyNet tries to overcome this problem with a **minimal mathematical documentation associated to each script** and a **wide range of comments inside the code**.
- ▶ Libraries like Tensorflow are efficient by a computational point-of-view and they have an extremely simple user interface. On the other hand the **deeper functionalities** of the code and the **implementation strategies** used are **unavoidably hidden behind tons of code lines**. In this way the user can performs complex computational tasks using the library as **black-box package**. NumPyNet wants to overcome this problem using **simple Python codes with extremely readability also for novel users** to better understand the symmetry between mathematical formulas and code.
- ▶ NumPyNet uses **Python threads** to easily manage independent tasks and optimize the computation.

\* The library was developed in collaboration with Dott. Ceccarelli.



# Byron

## Build YouR Own Neural Network

### Byron



- ▶ The library is written in **pure C++** with the support of the modern standard 17.
- ▶ The library is **optimized for image processing** (probably the most common task in biomedical research).
- ▶ Starting from the darknet project backbone, Byron proposes **numerous improvements and fixes**.
- ▶ Byron works in a **fully parallel section** (OpenMP) in which each single computational function is performed using the full set of available cores. To further reduce the time of thread spawn and so optimize as much as possible the code performances, the library **works using a single parallel section** which is opened at the beginning of the computation and closed at the end.
- ▶ The library is also **completely wrapped** using Cython to enlarge the range of users also to the Python ones.

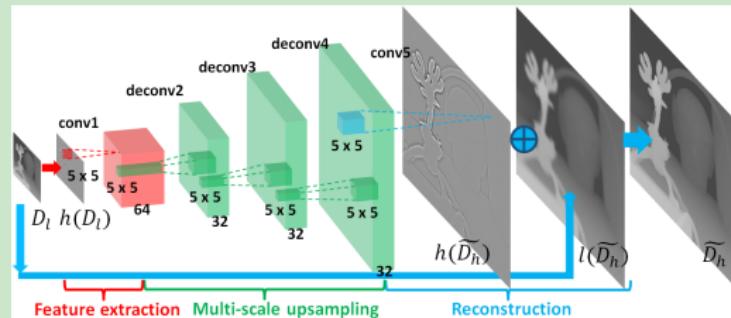
\* The library was developed in collaboration with Dott. Baroncini.



# Applications

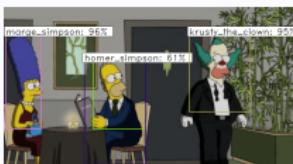
## Deep Learning model tasks

### Single Image Super Resolution



Improve spatial image resolution and reconstruction from low-resolution (LR) to high-resolution (HR).

### Object Detection



Identify single or multiple objects into a picture or video stream.

### Image Segmentation



Extract the exact pixels which belong to an object into a given image.



# Single Image Super Resolution

## What is Super Resolution?

We commonly think about Neural Network models as tools to **reduce** the problem dimensionality, i.e starting from high-dimensional data (e.g  $w \times h \times c$  image) the model predicts a class (single number).

In the Super Resolution problem we have no classes but the desired output is a image. Single Image Super Resolution tasks starts from a image aiming to reconstruct its HR counterpart.

### Super resolution procedure

- We start from a prior-known HR image;
- We rescale this image obtaining its LR counterpart;
- We feed a super-resolution model with the LR image aiming to obtain in output its HR version;
- With the trained model we can re-apply the up-sampling procedure and further increase the HR quality\*.

If this pipeline is performed feeding a Neural Network model using one-by-one a series of images, we talk about Single Image Super Resolution (SISR) algorithm.

\* This is even more efficient if we think about a **zoom** procedure.



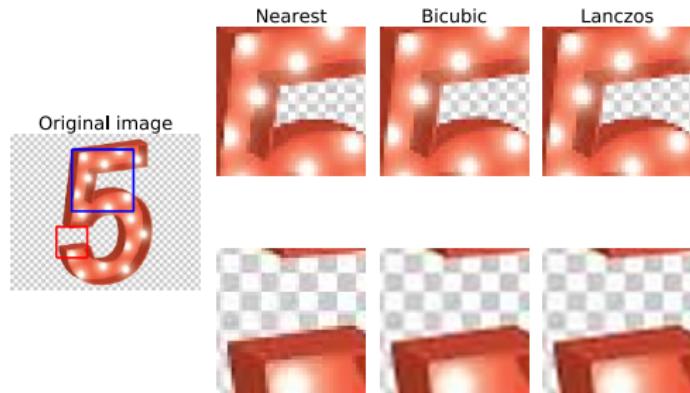
# Standard up/down scaling algorithms

## Single Image Super Resolution

Interpolation algorithms:

- Nearest;
- Bicubic;
- Lanczos;

The best compromise between computation cost and efficiency is given by the Bicubic interpolation algorithm. Given a pixel, the interpolation function evaluates the 4 pixels around it applying a filter given by the equation:



$$k(x) = \frac{1}{6} \begin{cases} (12 - 9B - 6C)|x|^3 + (-18 + 12B + 6C)|x|^2 + (6 - 2B) & \text{if } |x| < 1 \\ (-B - 6C)|x|^3 + (6B + 30C)|x|^2 + (-12B - 48C)|x| + (8B + 24C) & \text{if } 1 \leq |x| < 2 \\ 0 & \text{otherwise} \end{cases}$$

where  $x$  identifies each pixel below the filter. Commonly values used for the filter parameters are  $B = 0$  and  $C = 0.75$  (used by OpenCV library) or  $B = 0$  and  $C = 0.5$  used by Matlab.

**NOTE:** the nearest interpolation could be achieved also using *Pooling* algorithm



# Image Quality

## Single Image Super Resolution

The most common image quality evaluator is given by our eyes.

Quantitative evaluations:

### PSNR - Peak Signal to Noise Ratio

- ▶ It establishes the compression lossless of an image;
- ▶ Equation:

$$PSNR = 20 \cdot \log_{10} \left( \frac{\max(I)}{\sqrt{MSE}} \right)$$

where  $\max(I)$  is the maximum value which can be taken by a pixel in the image and  $MSE$  is the Mean Square Error evaluated between the original and reconstructed images.

### SSIM - Structural SIMilarity index

- ▶ It evaluates the structural similarity between two images taking into account also the visible improvement seen by human eyes.
- ▶ Equation:

$$SSIM(I, K) = \frac{1}{N} \sum_{i=1}^N \frac{(2\mu_x\mu_y + c_1)(2\sigma_{xy} + c_2)}{(\mu_x^2 + \mu_y^2 + c_1)(\sigma_x^2 + \sigma_y^2 + c_2)}$$

where  $N$  is the number of arbitrary patches which divide the image and  $c_1$  and  $c_2$  parameters are fixed to avoid mathematical divergence.

	PSNR (dB)	SSIM ( $\in [-1, 1]$ )
Nearest	25.118	0.847
Bicubic	27.254	0.894
Lanczos	26.566	0.871



# EDSR - Enhanced Deep Super Resolution

## Super Resolution Models

Immagine HR  $\rightarrow$  ricampionamento bicubico  $\rightarrow$  Immagine LR  $\rightarrow$  super-risoluzione  $\rightarrow$  Immagine SR



### Architecture

Layer	Channels input/output	Filter dimensions	Number of Parameters
Conv. input	3/256	$3 \times 3$	6912
Conv. (residual block)	256/256	$3 \times 3$	589824
conv. (pre-shuffle)	256/256	$3 \times 3$	589824
Conv. (upsample block)	256/1024	$3 \times 3$	2359296
Conv. output	256/3	$3 \times 3$	6912

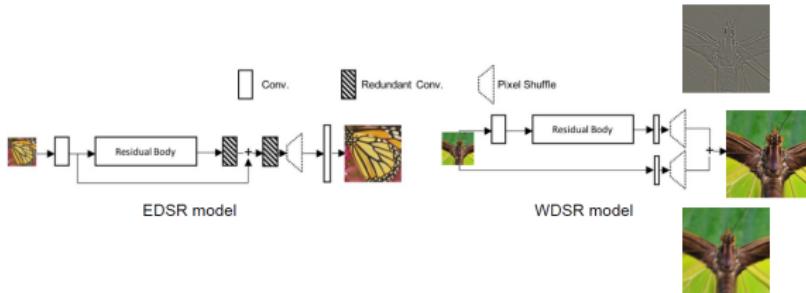
Average Time on  $510 \times 339$  image: 576.92 s.

Winner of the **NTIRE 2017**. More than **3 million** of parameters!



# WDSR - Wide Deep Super Resolution

## Super Resolution Models



### Architecture

Layer	Channels input/output	Filter dimensions	Number of Parameters
Conv. input 1	3/32	$3 \times 3$	864
Conv. 1 (residual block)	32/192	$3 \times 3$	55296
conv. 2 (residual block)	192/32	$3 \times 3$	55296
Conv. (pre-shuffle)	32/48	$3 \times 3$	13824
Conv. input 2 (pre-shuffle)	3/48	$5 \times 5$	3600

Average Time on  $510 \times 339$  image: 46.35 s.

Winner of the NTIRE 2018.  $\sim 100K$  parameters, less than 10% of EDSR model!

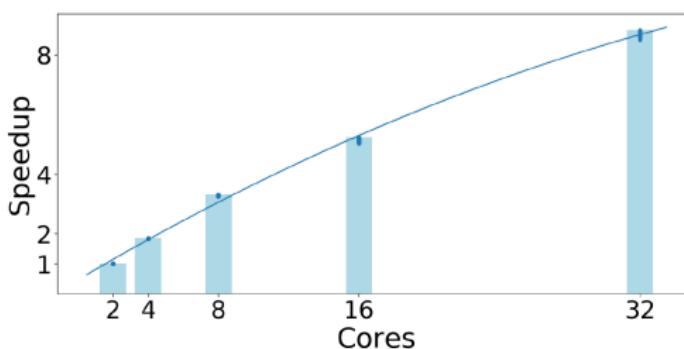


# Computational Performances

## Super Resolution Models

Computational performances of WDSR model on  $510 \times 339$  images (100 runs).

Environment: Bioinformatics server grade machine (128 GB RAM memory and 2 CPU E5-2620, with 8 cores each).



N° threads	Average Time (s)
2	403.478
4	219.268
8	122.986
16	78.571
32	46.348

**Sub-linear trend:**

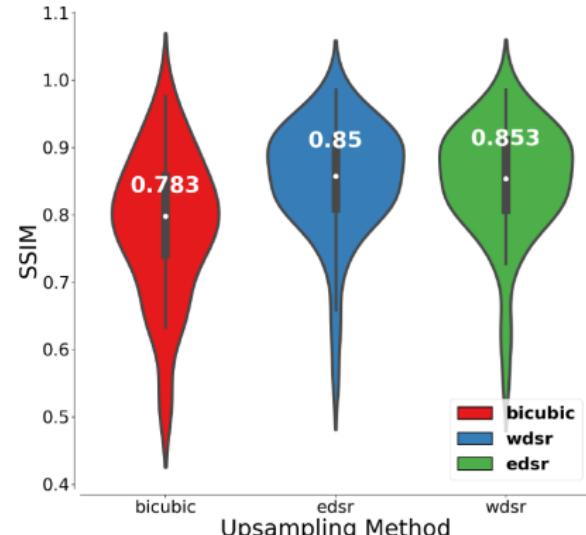
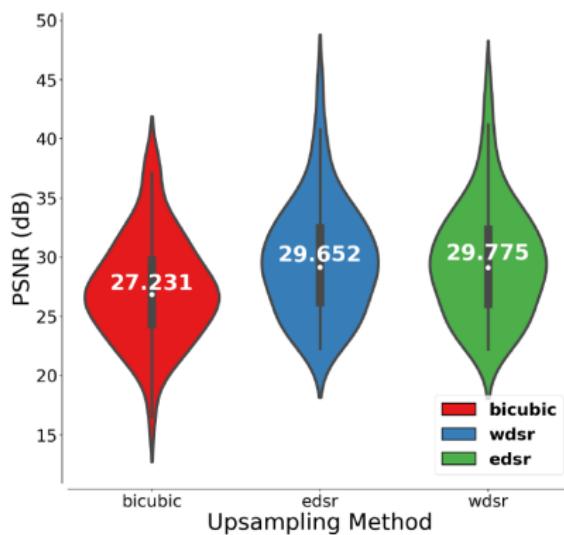
- Hyperthreading;
- NO affinity;



# Quality Performances

## Super Resolution Models

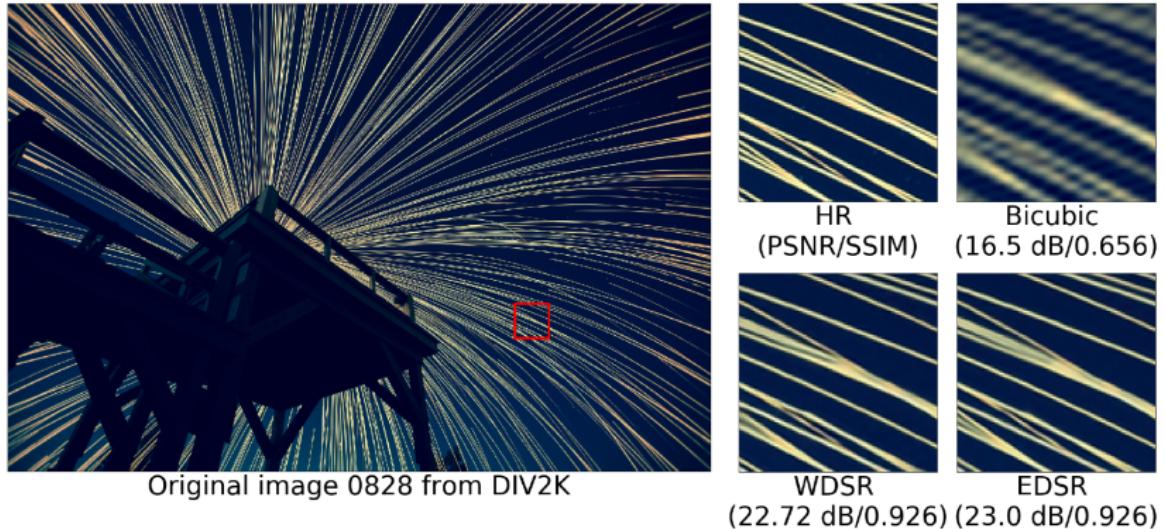
Evaluation of PSNR and SSIM on DIV2K validation set.





# Image Results

## Single Image Super Resolution



The model is trained on the **DIV2K** (*DIVerse 2K resolution high quality images*) dataset.

The dataset contains 800 high-resolution images as training set and their corresponding low-resolution ones, obtained by different down-sampling methods and different scale factors (2, 3, and 4).



# Image Results

## Single Image Super Resolution



Original image 0861 from DIV2K



HR  
(PSNR/SSIM)



Bicubic  
(21.23 dB/0.648)



WDSR  
(23.17 dB/0.776)



EDSR  
(23.36 dB/0.783)

The model is trained on the **DIV2K** (*DIVerse 2K resolution high quality images*) dataset.

The dataset contains 800 high-resolution images as training set and their corresponding low-resolution ones, obtained by different down-sampling methods and different scale factors (2, 3, and 4).

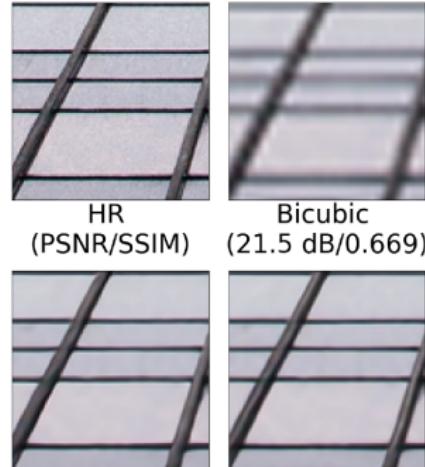


# Image Results

## Single Image Super Resolution



Original image 0845 from DIV2K



HR  
(PSNR/SSIM)

Bicubic  
(21.5 dB/0.669)

WDSR

(24.58 dB/0.801)

EDSR

(25.04 dB/0.811)

The model is trained on the **DIV2K** (*DIVerse 2K resolution high quality images*) dataset.

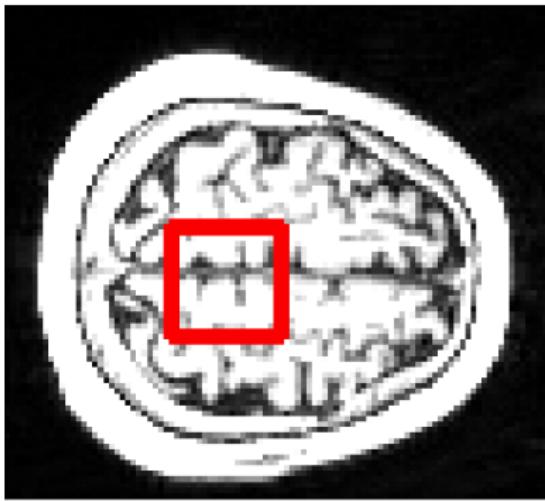
The dataset contains 800 high-resolution images as training set and their corresponding low-resolution ones, obtained by different down-sampling methods and different scale factors (2, 3, and 4).



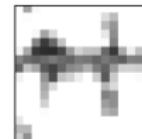
# Model Extrapolation

## NMR Application

Test on T1 weighted NMR images ( $256 \times 256 \rightarrow \text{REAL TIME}$ ).  
The images were down-sampled ( $128 \times 128 \rightarrow 2x$  and  $64 \times 64 \rightarrow 4x$ ) and then  
re-up-sampled.



Raw ROI



Bicubic Upscale



Super Resolution



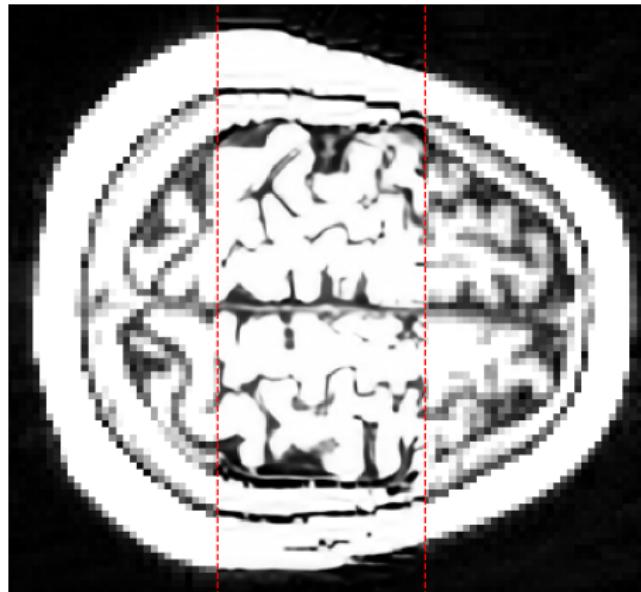


# Model Extrapolation

## NMR Application

Test on T1 weighted NMR images ( $256 \times 256 \rightarrow \text{REAL TIME}$ ).

The images were down-sampled ( $128 \times 128 \rightarrow 2x$  and  $64 \times 64 \rightarrow 4x$ ) and then re-up-sampled.





# Visual score

## NMR Application

---

Test on T1 weighted NMR images ( $256 \times 256 \rightarrow \text{REAL TIME}$ ).  
The images were down-sampled ( $128 \times 128 \rightarrow 2x$  and  $64 \times 64 \rightarrow 4x$ ) and then  
re-up-sampled.

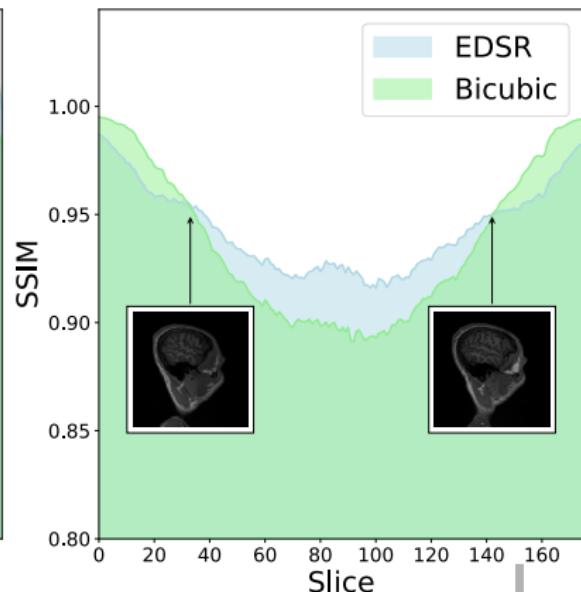
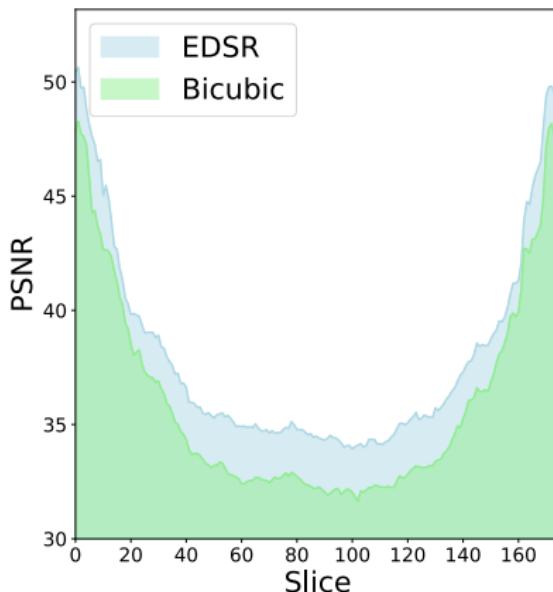


# Quality score

## NMR Application

Test on T1 weighted NMR images ( $256 \times 256 \rightarrow$  **REAL TIME**).

The images were down-sampled ( $128 \times 128 \rightarrow 2x$  and  $64 \times 64 \rightarrow 4x$ ) and then re-up-sampled with SR models.



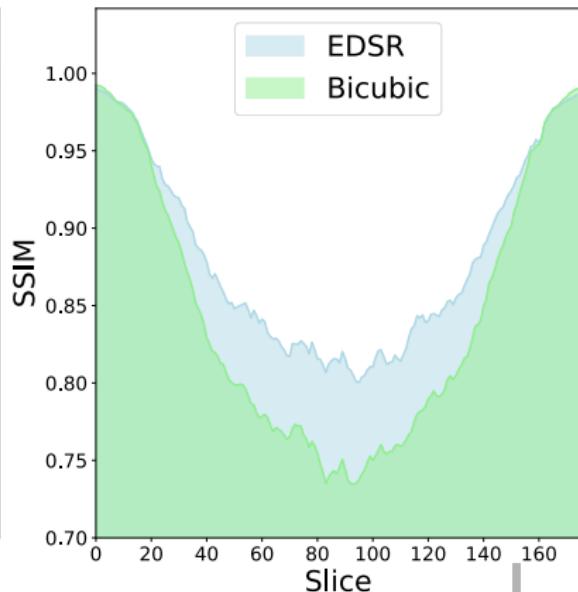
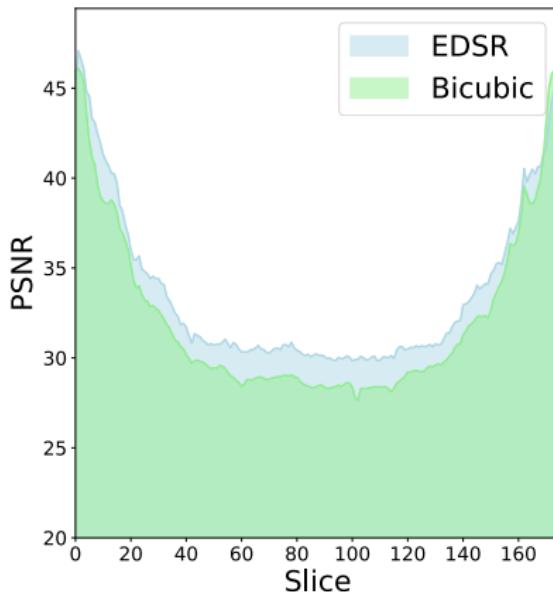


# Quality score

## NMR Application

Test on T1 weighted NMR images ( $256 \times 256 \rightarrow$  **REAL TIME**).

The images were down-sampled ( $128 \times 128 \rightarrow 2x$  and  $64 \times 64 \rightarrow 4x$ ) and then re-up-sampled with SR models.

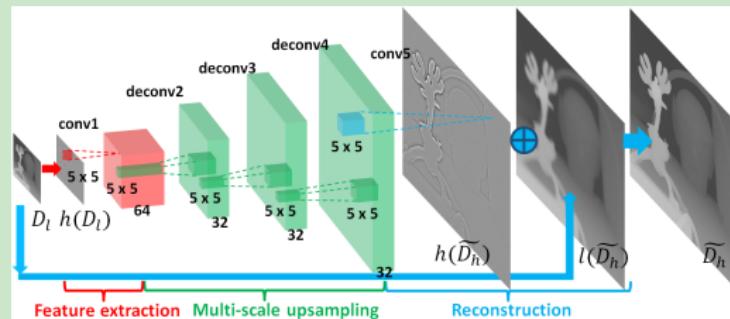




# Applications

## Deep Learning model tasks

### Single Image Super Resolution



Improve spatial image resolution and reconstruction from low-resolution (LR) to high-resolution (HR).

### Object Detection



Identify single or multiple objects into a picture or video stream.

### Image Segmentation



Extract the exact pixels which belong to an object into a given image.



# Object Detection

## What is Object Detection?

Object detection is one of the larger deep learning sub-discipline, especially when we talk about Neural Network models.

This kind of problems aim to identify single or multiple objects into a picture or video stream.

Possible applications:

- **Object tracking;**
- Video surveillance;
- **Pedestrian detection;**
- Anomaly detection;
- **People counting;**
- Self-driving cars;
- Face detection;
- ...;



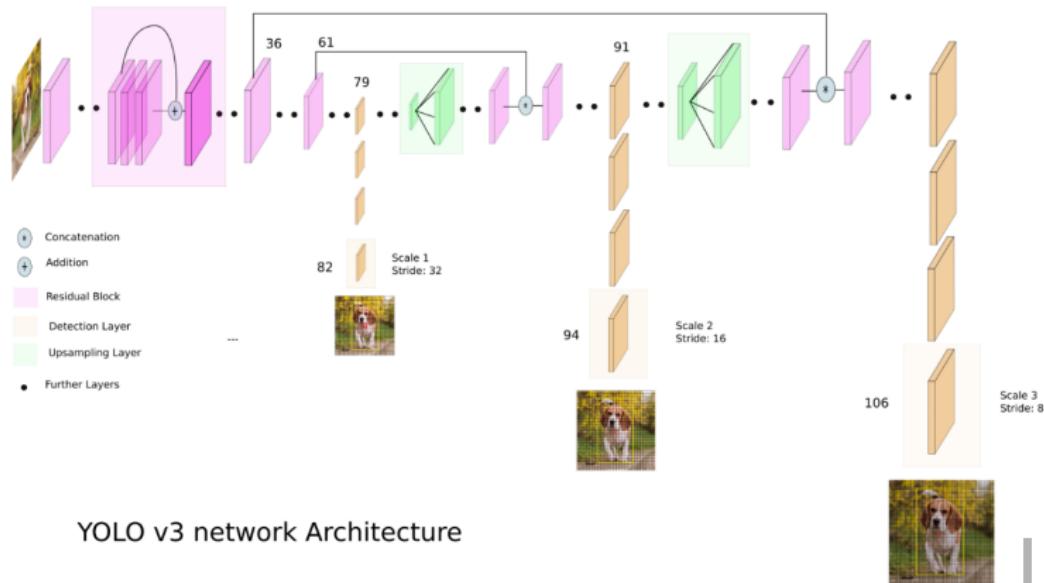


# YOLO - You Only Look Once

## Object Detection Model

YOLO is a deep Neural Network model with more than 100 layers and more than 62 million of parameters.

The original implementation of the model is provided by Joseph Redmon (<https://github.com/pjreddie/darknet>) into the **darknet** project (ANSI C).



YOLO v3 network Architecture

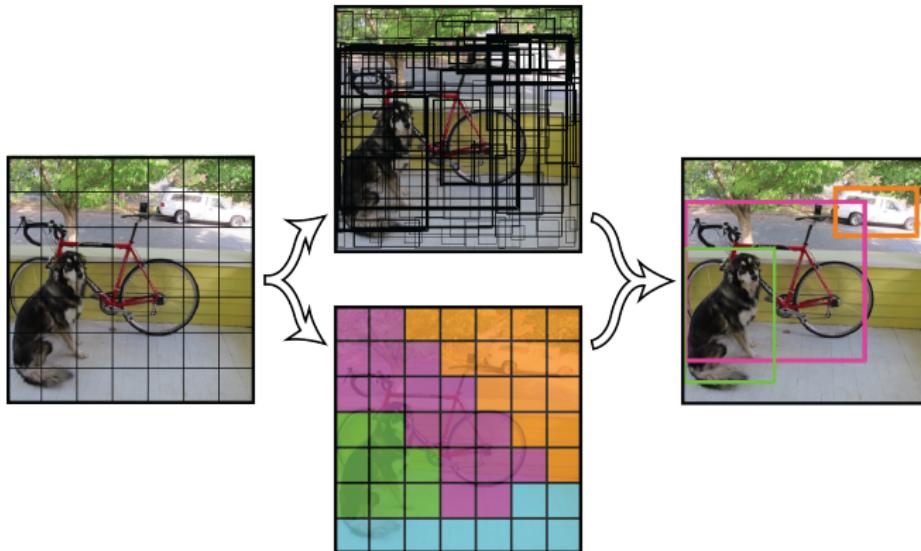


# YOLO - You Only Look Once

## Object Detection Model

YOLO is a deep Neural Network model with more than 100 layers and more than 62 million of parameters.

The original implementation of the model is provided by Joseph Redmon (<https://github.com/pjreddie/darknet>) into the **darknet** project (ANSI C).





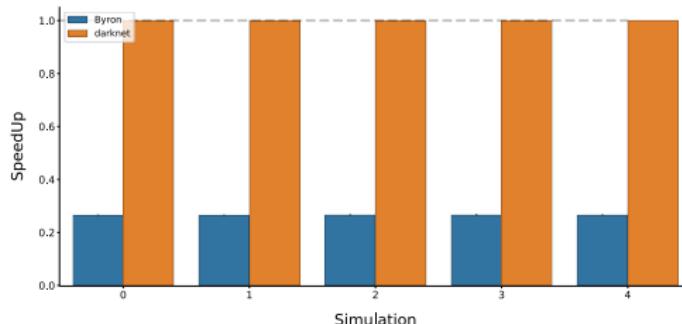
# Computational Performances

## Object Detection Models

Byron is inspired to darknet and its first implementation was about object detection task.

darknet → benchmark → but its efficiency is based on GPU support!

Evaluation performed on 100 runs without multi-threading!



Framework	Average Time (s)
Byron	5.18
darknet	19.58

Time evaluated on  $416 \times 416$  images.

Environment: Bioinformatics server grade machine (128 GB RAM memory and 2 CPU E5-2620, with 8 cores each).



# Model Extrapolation

## People counting Application

Despite the model is incredibly efficient in object detection also with low quality images, there is a sort of **limit** in the number of pixels needed for object identification. We had the opportunity to empirically verify its limit working on a people tracking project for real time applications.

The project was developed in collaboration with the Complex Systems (**PhySyCom**) group.

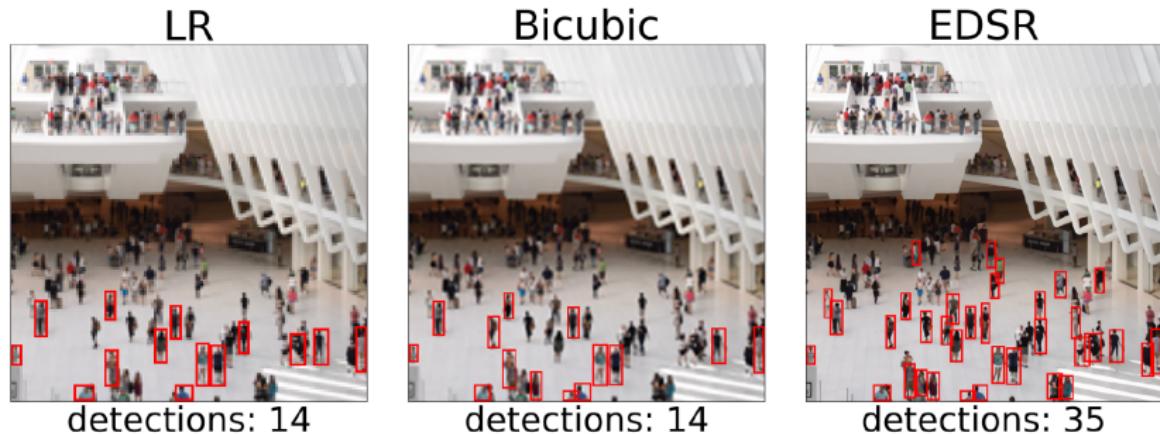


# Model Extrapolation

## People counting Application

Despite the model is incredibly efficient in object detection also with low quality images, there is a sort of **limit** in the number of pixels needed for object identification. We had the opportunity to empirically verify its limit working on a people tracking project for real time applications.

The project was developed in collaboration with the Complex Systems (**PhySyCom**) group.





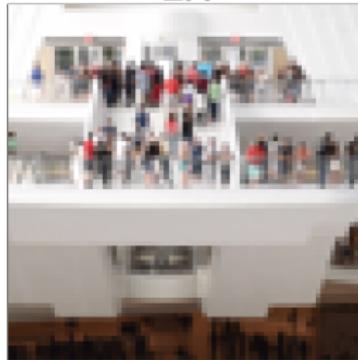
# Model Extrapolation

## People counting Application

Despite the model is incredibly efficient in object detection also with low quality images, there is a sort of **limit** in the number of pixels needed for object identification. We had the opportunity to empirically verify its limit working on a people tracking project for real time applications.

The project was developed in collaboration with the Complex Systems (**PhySyCom**) group.

LR



Bicubic



EDSR



detections: 0

detections: 2

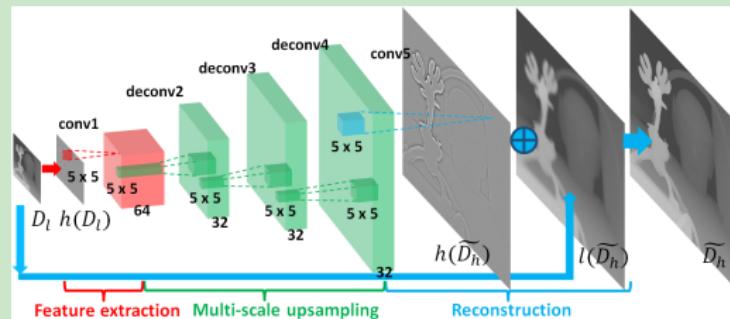
detections: 7



# Applications

## Deep Learning model tasks

### Single Image Super Resolution



Improve spatial image resolution and reconstruction from low-resolution (LR) to high-resolution (HR).

### Object Detection



Identify single or multiple objects into a picture or video stream.

### Image Segmentation



Extract the exact pixels which belong to an object into a given image.



# Image Segmentation

## What is Image Segmentation?

**Object Detection:** give a label to **ROIs** of the input image.

**Image Segmentation:** give a label to **each pixel** of the input image.





# Image Segmentation in Medicine

## 3D Femur Structure

### Rizzoli Hospital project

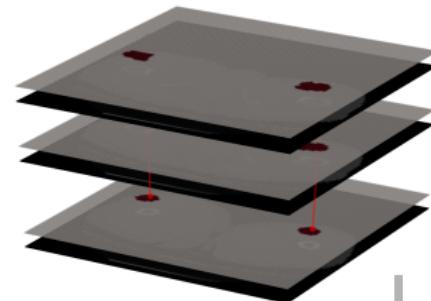
Improve the identification and segmentation of the femoral head, trying to discriminate this part of the bone from the articular cartilage and moreover from the acetabular fossa.

**Data type:** CT (*Computer Tomography*) images.

**Dataset:** 4 (not-annotated) patients.

### Semi-automatic pipeline:

- Edge enhancement;
- Thresholding;
- Morphological operations;
- Connected Components;
- Interpolation between consecutive frames;

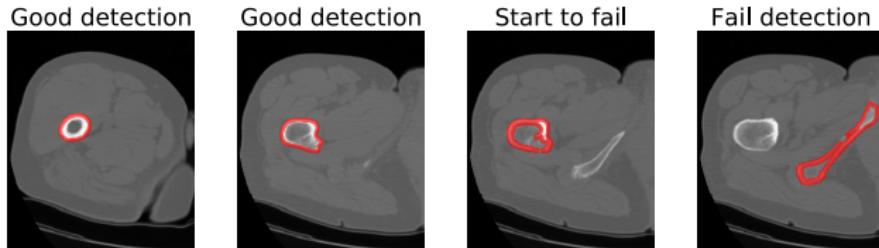




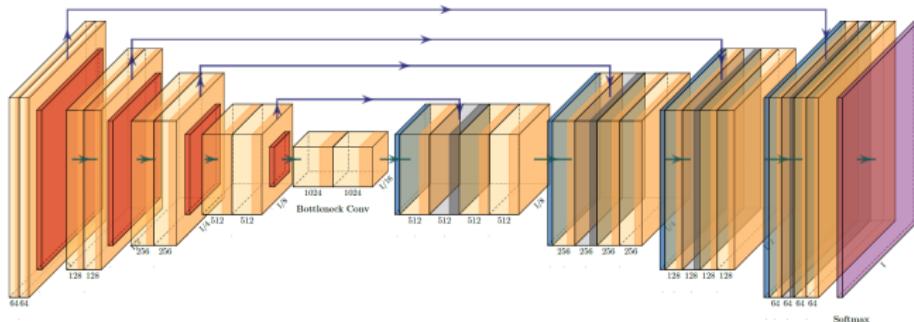
# Image Segmentation in Medicine

## Segmentation Models

The method is too naive to perform a good segmentation on the full set of slices.  
However, it can be useful to **reduce** the quantity of slices to annotate manually ( $\sim 400$  slices per patient)!



Thus, we moved to a **Neural Network model**





# U-Net Model

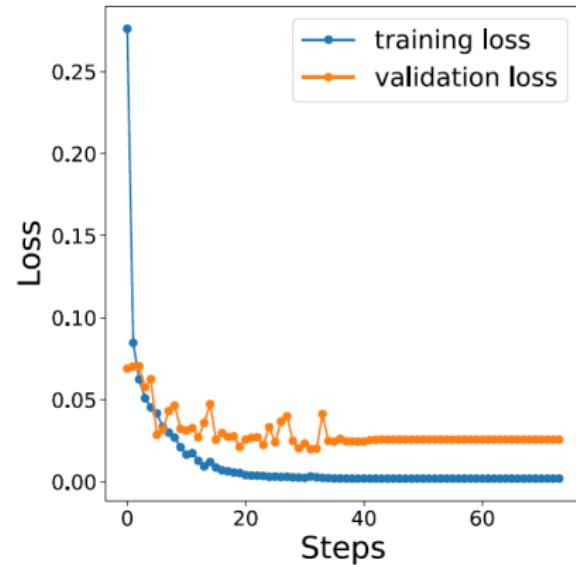
## Segmentation Models

Training parameters:

- Training Set: 96 images;
- Test Set: 40 images;
- Data augmentation: Rotated, shifted, mirrored.
- 40 epochs;
- loss: binary cross-entropy
- metrics: IoU (*Intersection over Union*)

### IoU Validation Set:

The 80% of the test set has obtained a IoU score greater than 0.8 and thus a good correspondence between our results and the ground truth.





# IoU score

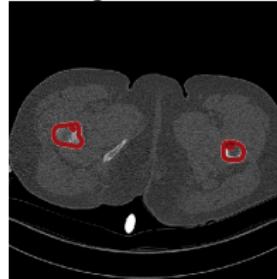
## Segmentation Models

The model output were filtered using a thresholding of  $10^{-2}$ , i.e values  $\leq$  were turned off.

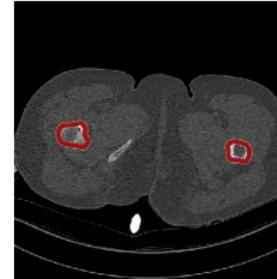
Output mask



Segmentation

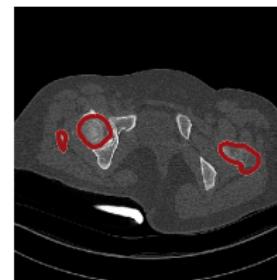


Ground Truth



IoU Score

0.881



0.801



# 3D Reconstruction

## Segmentation Models



# Image Segmentation in Medicine

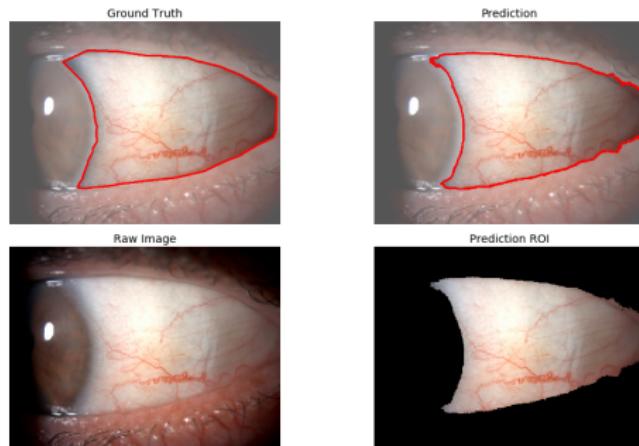
## ROSSO project

### ROSSO - Robotic Ocular Semantic Segmentation for Ophthalmology

Sant'Orsola Hospital project: Extract the pixels belonging to sclera for the quantification of the red part of the image.

All the images have been manually annotated by experts. Dataset:

- Training: 150 images;
- Test: 18 images;





# Image Segmentation in Medicine

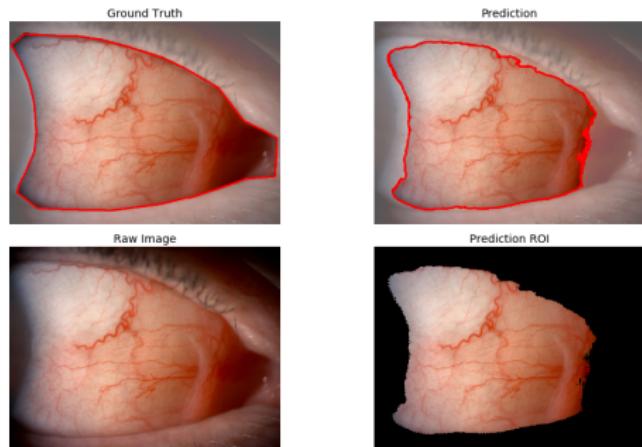
## ROSSO project

### ROSSO - Robotic Ocular Semantic Segmentation for Ophthalmology

Sant'Orsola Hospital project: Extract the pixels belonging to sclera for the quantification of the red part of the image.

All the images have been manually annotated by experts. Dataset:

- Training: 150 images;
- Test: 18 images;





# Image Segmentation in Medicine

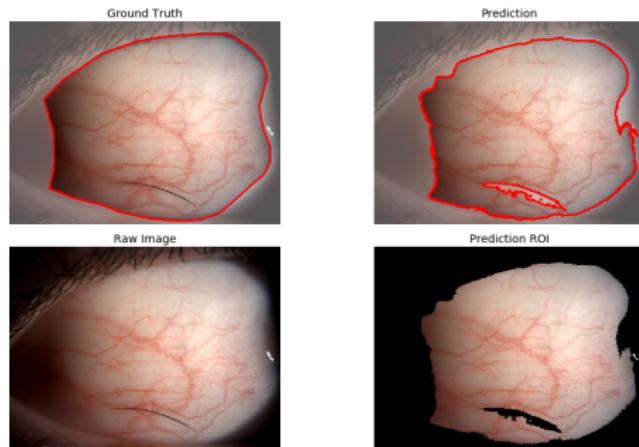
## ROSSO project

### ROSSO - Robotic Ocular Semantic Segmentation for Ophthalmology

Sant'Orsola Hospital project: Extract the pixels belonging to sclera for the quantification of the red part of the image.

All the images have been manually annotated by experts. Dataset:

- Training: 150 images;
- Test: 18 images;



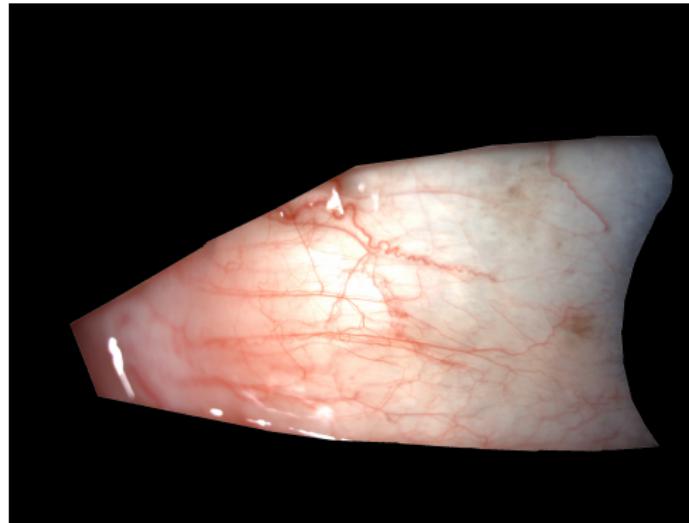


# Image Segmentation in Medicine

## ROSSO project

### ROSSO - Robotic Ocular Semantic Segmentation for Ophthalmology

Sant'Orsola Hospital project: Extract the vascular network from the eye's sclera component.



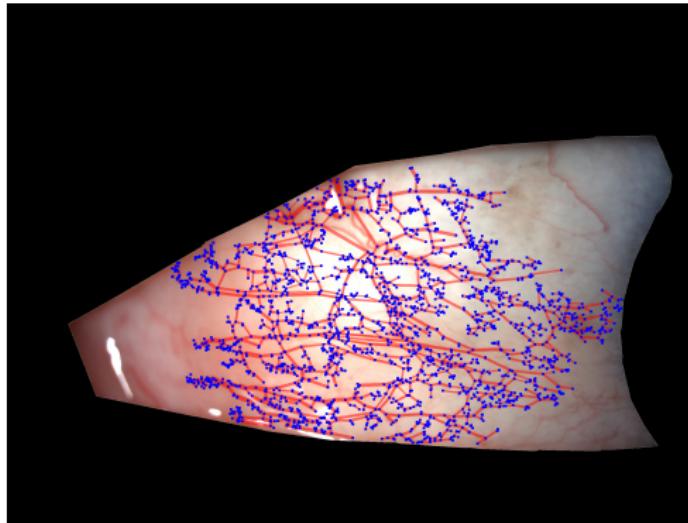


# Image Segmentation in Medicine

## ROSSO project

### ROSSO - Robotic Ocular Semantic Segmentation for Ophthalmology

Sant'Orsola Hospital project: Extract the vascular network from the eye's sclera component.





# Conclusion

- Two new libraries for Deep Learning applications were proposed:
  - **NumPyNet**: a tool for educational practice and modeling;
  - **Byron**: a tool for an efficient implementation of these models;
- Both libraries implement cutting edge numerical techniques to maximize computation efficiency;
- **Byron** optimize and extend the **darknet** project with several fixes and computational improvements.
- **Byron** is able to join in a single framework either super resolution and object detection tasks.
- The results obtained using Super Resolution and Segmentation models are very promising for the analysis of biomedical images.
- We have shown that deep learning models are capable of a very efficient generalization due to their vast amount of parameters and a well-programmed training section.
- The **Byron** YOLO version is approximately 3.8x faster than **darknet** in all the simulations;



# Project References

<https://github.com/Nico-Curti>

- 
- **Byron**: Build YouR Own Neural Network, <http://github.com/Nico-Curti/Byron>
  - **NumPyNet**: Neural Network library in Pure Numpy, <http://github.com/Nico-Curti/NumPyNet>
  - **DNetPRO**: Discriminant Analysis with Network PROcessing, <http://github.com/Nico-Curti/DNetPRO>
  - **rFBP**: Replicated Focusing Belief Propagation algorithm, <http://github.com/Nico-Curti/rFBP>
  - **Scorer**: Machine Learning Scorer Library with parallel DAG support, <http://github.com/Nico-Curti/scorer>
  - **FiloBluService**: FiloBlu project service manager for text message processing, <http://github.com/Nico-Curti/FiloBluService>
  - **ShUt**: Shell Utilities and Installers for no root users, <http://github.com/Nico-Curti/shut>
  - **genetic**: Examples about genetic algorithms for parallel and distributed computing, <http://github.com/Nico-Curti/genetic>
  - **BlendNet**: Network viewer with Blender support, <http://github.com/Nico-Curti/BlendNet>
  - **CryptoSocket**: TCP/IP Client Server with RSA cryptography, <http://github.com/Nico-Curti/CryptoSocket>
  - **easyDAG**: Simple template DAG scheduler in C++, <http://github.com/Nico-Curti/easyDAG>
  - **Cardio**: Pulse oximetry data processing and classification, <http://github.com/Nico-Curti/cardio>
  - **SysDyn**: System Dynamics script utilities, <http://github.com/Nico-Curti/SysDyn>
  - **rSGD**: Replicated Stochastic Gradient Descent algorithm, <http://github.com/Nico-Curti/rSGD>
  - **Walkers**: Random Walk and Optimizer Simulator, <http://github.com/Nico-Curti/Walkers>
  - **Data-Analysis**: Data Analysis Utilities (from Statistics to Machine Learning), <http://github.com/Nico-Curti/Data-ANalysis>



# Thanks to the research group

